

## **SQL Server 2008 Query Performance Tuning Distilled**

**Copyright © 2009 by Grant Fritchey and Sajal Dam**

All rights reserved. No part of this work may be reproduced or transmitted in any form or by any means, electronic or mechanical, including photocopying, recording, or by any information storage or retrieval system, without the prior written permission of the copyright owner and the publisher.

ISBN-13 (pbk): 978-1-4302-1902-6

ISBN-13 (electronic): 978-1-4302-1903-3

Printed and bound in the United States of America 9 8 7 6 5 4 3 2 1

Trademarked names may appear in this book. Rather than use a trademark symbol with every occurrence of a trademarked name, we use the names only in an editorial fashion and to the benefit of the trademark owner, with no intention of infringement of the trademark.

Lead Editor: Jonathan Gennick

Development Editor: Douglas Pundick

Technical Reviewer: Joseph Sack

Editorial Board: Clay Andres, Steve Anglin, Mark Beckner, Ewan Buckingham, Tony Campbell,

Gary Cornell, Jonathan Gennick, Michelle Lowman, Matthew Moodie, Jeffrey Pepper, Frank Pohlmann,

Ben Renow-Clarke, Dominic Shakeshaft, Matt Wade, Tom Welsh

Project Manager: Richard Dal Porto

Copy Editor: Kim Wimpsett

Associate Production Director: Kari Brooks-Copony

Production Editor: Kelly Winquist

Compositor: Patrick Cunningham

Proofreader: April Eddy

Indexer: John Collin

Artist: April Milne

Cover Designer: Kurt Krames

Manufacturing Director: Tom Debolski

Distributed to the book trade worldwide by Springer-Verlag New York, Inc., 233 Spring Street, 6th Floor, New York, NY 10013. Phone 1-800-SPRINGER, fax 201-348-4505, e-mail [orders-ny@springer-sbm.com](mailto:orders-ny@springer-sbm.com), or visit <http://www.springeronline.com>.

For information on translations, please contact Apress directly at 2855 Telegraph Avenue, Suite 600, Berkeley, CA 94705. Phone 510-549-5930, fax 510-549-5939, e-mail [info@apress.com](mailto:info@apress.com), or visit <http://www.apress.com>.

Apress and friends of ED books may be purchased in bulk for academic, corporate, or promotional use. eBook versions and licenses are also available for most titles. For more information, reference our Special Bulk Sales—eBook Licensing web page at <http://www.apress.com/info/bulksales>.

The information in this book is distributed on an “as is” basis, without warranty. Although every precaution has been taken in the preparation of this work, neither the author(s) nor Apress shall have any liability to any person or entity with respect to any loss or damage caused or alleged to be caused directly or indirectly by the information contained in this work.

The source code for this book is available to readers at <http://www.apress.com>.

# Contents

About the Author .....	xix
About the Technical Reviewer .....	xxi
Acknowledgments .....	xxiii
Introduction .....	xxv

<b>CHAPTER 1</b>	<b>SQL Query Performance Tuning .....</b>	<b>1</b>
	The Performance-Tuning Process .....	2
	The Core Process .....	2
	Iterating the Process .....	4
	Performance vs. Price .....	7
	Performance Targets .....	7
	“Good Enough” Tuning .....	7
	Performance Baseline .....	8
	Where to Focus Efforts .....	9
	SQL Server Performance Killers .....	10
	Poor Indexing .....	11
	Inaccurate Statistics .....	11
	Excessive Blocking and Deadlocks .....	11
	Non-Set-Based Operations .....	12
	Poor Query Design .....	12
	Poor Database Design .....	12
	Excessive Fragmentation .....	13
	Nonreusable Execution Plans .....	13
	Poor Execution Plans .....	13
	Frequent Recompilation of Execution Plans .....	14
	Improper Use of Cursors .....	14
	Improper Configuration of the Database Log .....	14
	Excessive Use or Improper Configuration of tempdb .....	14
	Summary .....	15

<b>CHAPTER 2</b>	<b>System Performance Analysis</b>	17
	Performance Monitor Tool	17
	Dynamic Management Views	19
	Hardware Resource Bottlenecks	20
	Identifying Bottlenecks	20
	Bottleneck Resolution	21
	Memory Bottleneck Analysis	21
	SQL Server Memory Management	22
	Available Bytes	25
	Pages/sec and Page Faults/sec Counters	25
	Buffer Cache Hit Ratio	26
	Page Life Expectancy	26
	Checkpoint Pages/sec	27
	Lazy writes/sec	27
	Memory Grants Pending	27
	Target Server Memory (KB) and Total Server Memory (KB)	27
	Memory Bottleneck Resolutions	27
	Optimizing Application Workload	29
	Allocating More Memory to SQL Server	29
	Increasing System Memory	29
	Changing from a 32-bit to a 64-bit Processor	29
	Enabling 3GB of Process Space	30
	Using Memory Beyond 4GB Within SQL Server	30
	Disk Bottleneck Analysis	32
	Disk Counters	32
	% Disk Time	33
	Current Disk Queue Length	33
	Disk Transfers/sec	34
	Disk Bytes/sec	34
	Avg. Disk Sec/Read and Avg. Disk Sec/Write	34
	Disk Bottleneck Resolutions	35
	Optimizing Application Workload	35
	Using a Faster Disk Drive	35
	Using a RAID Array	35
	Using a SAN System	37
	Aligning Disks Properly	38
	Using a Battery-Backed Controller Cache	38
	Adding System Memory	38

Creating Multiple Files and Filegroups . . . . .	39
Placing the Table and Index on Separate Disks . . . . .	42
Saving Log Files to a Separate Physical Disk . . . . .	42
Partitioning Tables . . . . .	43
Processor Bottleneck Analysis . . . . .	43
% Processor Time . . . . .	44
% Privileged Time . . . . .	44
Processor Queue Length . . . . .	44
Context Switches/sec . . . . .	44
Batch Requests/sec . . . . .	45
SQL Compilations/sec . . . . .	45
SQL Recompilations/sec . . . . .	45
Processor Bottleneck Resolutions . . . . .	45
Optimizing Application Workload . . . . .	45
Eliminating Excessive Compiles/Recompiles . . . . .	46
Using More or Faster Processors . . . . .	46
Using a Large L2/L3 Cache . . . . .	46
Running More Efficient Controllers/Drivers . . . . .	47
Not Running Unnecessary Software . . . . .	47
Network Bottleneck Analysis . . . . .	47
Bytes Total/sec . . . . .	48
% Net Utilization . . . . .	48
Network Bottleneck Resolutions . . . . .	48
Optimizing Application Workload . . . . .	48
Adding Network Adapters . . . . .	49
Moderating and Avoiding Interruptions . . . . .	49
SQL Server Overall Performance . . . . .	49
Missing Indexes . . . . .	50
Database Blocking . . . . .	51
Nonreusable Execution Plans . . . . .	52
General Behavior . . . . .	52
Creating a Baseline . . . . .	53
Creating a Reusable List of Performance Counters . . . . .	53
Creating a Counter Log Using the List of Performance Counters . . . . .	56
Minimizing Performance Monitor Overhead . . . . .	57
System Behavior Analysis Against Baseline . . . . .	59
Summary . . . . .	60

<b>CHAPTER 3</b>	<b>SQL Query Performance Analysis</b>	61
	The SQL Profiler Tool	61
	Profiler Traces	62
	Events	63
	Data Columns	66
	Filters	67
	Trace Templates	69
	Trace Data	69
	Trace Automation	70
	Capturing a Trace Using the GUI	70
	Capturing a Trace Using Stored Procedures	71
	Combining Trace and Performance Monitor Output	72
	SQL Profiler Recommendations	73
	Limiting the Number of Events and Data Columns	73
	Discarding Start Events for Performance Analysis	74
	Limiting the Trace Output Size	75
	Avoiding Online Data Column Sorting	75
	Running Profiler Remotely	75
	Limiting the Use of Certain Events	75
	Query Performance Metrics Without Profiler	76
	Costly Queries	76
	Identifying Costly Queries	77
	Identifying Slow-Running Queries	82
	Execution Plans	83
	Analyzing a Query Execution Plan	84
	Identifying the Costly Steps in an Execution Plan	87
	Analyzing Index Effectiveness	87
	Analyzing Join Effectiveness	89
	Actual vs. Estimated Execution Plans	93
	Plan Cache	95
	Query Cost	95
	Client Statistics	96
	Execution Time	97
	STATISTICS IO	98
	Summary	100
<b>CHAPTER 4</b>	<b>Index Analysis</b>	101
	What Is an Index?	101
	The Benefit of Indexes	103
	Index Overhead	105

Index Design Recommendations . . . . .	107
Examine the WHERE Clause and Join Criteria Columns . . . . .	107
Use Narrow Indexes . . . . .	109
Examine Column Uniqueness . . . . .	111
Examine the Column Data Type . . . . .	114
Consider Column Order . . . . .	114
Consider the Type of Index . . . . .	117
Clustered Indexes . . . . .	117
Heap Tables . . . . .	118
Relationship with Nonclustered Indexes . . . . .	118
Clustered Index Recommendations . . . . .	120
Nonclustered Indexes . . . . .	126
Nonclustered Index Maintenance . . . . .	127
Defining the Bookmark Lookup . . . . .	127
Nonclustered Index Recommendations . . . . .	127
Clustered vs. Nonclustered Indexes . . . . .	128
Benefits of a Clustered Index over a Nonclustered Index . . . . .	129
Benefits of a Nonclustered Index over a Clustered Index . . . . .	131
Advanced Indexing Techniques . . . . .	132
Covering Indexes . . . . .	132
Index Intersections . . . . .	135
Index Joins . . . . .	137
Filtered Indexes . . . . .	138
Indexed Views . . . . .	141
Index Compression . . . . .	145
Special Index Types . . . . .	147
Full-Text . . . . .	147
Spatial . . . . .	148
XML . . . . .	148
Additional Characteristics of Indexes . . . . .	148
Different Column Sort Order . . . . .	148
Index on Computed Columns . . . . .	148
Index on BIT Data Type Columns . . . . .	149
CREATE INDEX Statement Processed As a Query . . . . .	149
Parallel Index Creation . . . . .	149
Online Index Creation . . . . .	150
Considering the Database Engine Tuning Advisor . . . . .	150
Summary . . . . .	150

<b>CHAPTER 5</b>	<b>Database Engine Tuning Advisor</b>	151
	Database Engine Tuning Advisor Mechanisms	151
	Database Engine Tuning Advisor Examples	155
	Tuning a Query	155
	Tuning a Trace Workload	159
	Database Engine Tuning Advisor Limitations	161
	Summary	162
 <b>CHAPTER 6</b>	 <b>Bookmark Lookup Analysis</b>	 163
	Purpose of Bookmark Lookups	163
	Drawbacks of Bookmark Lookups	165
	Analyzing the Cause of a Bookmark Lookup	166
	Resolving Bookmark Lookups	169
	Using a Clustered Index	169
	Using a Covering Index	169
	Using an Index Join	173
	Summary	174
 <b>CHAPTER 7</b>	 <b>Statistics Analysis</b>	 175
	The Role of Statistics in Query Optimization	175
	Statistics on an Indexed Column	176
	Benefits of Updated Statistics	177
	Drawbacks of Outdated Statistics	179
	Statistics on a Nonindexed Column	180
	Benefits of Statistics on a Nonindexed Column	181
	Drawback of Missing Statistics on a Nonindexed Column	185
	Analyzing Statistics	187
	Density	190
	Statistics on a Multicolumn Index	191
	Statistics on a Filtered Index	192
	Statistics Maintenance	193
	Automatic Maintenance	193
	Manual Maintenance	195
	Statistics Maintenance Status	198
	Analyzing the Effectiveness of Statistics for a Query	199
	Resolving a Missing Statistics Issue	199
	Resolving an Outdated Statistics Issue	202

Recommendations . . . . .	204
Backward Compatibility of Statistics . . . . .	204
Auto Create Statistics . . . . .	204
Auto Update Statistics . . . . .	205
Automatic Update Statistics Asynchronously . . . . .	207
Amount of Sampling to Collect Statistics . . . . .	208
Summary . . . . .	208
 <b>CHAPTER 8    Fragmentation Analysis . . . . .</b>	 <b>209</b>
Causes of Fragmentation . . . . .	209
Page Split by an UPDATE Statement . . . . .	212
Page Split by an INSERT Statement . . . . .	215
Fragmentation Overhead . . . . .	217
Analyzing the Amount of Fragmentation . . . . .	220
Analyzing the Fragmentation of a Small Table . . . . .	222
Fragmentation Resolutions . . . . .	224
Dropping and Re-creating the Index . . . . .	225
Re-creating the Index with the DROP_EXISTING Clause . . . . .	225
Executing the ALTER INDEX REBUILD Statement . . . . .	226
Executing the ALTER INDEX REORGANIZE Statement . . . . .	228
Significance of the Fill Factor . . . . .	230
Automatic Maintenance . . . . .	233
Summary . . . . .	239
 <b>CHAPTER 9    Execution Plan Cache Analysis . . . . .</b>	 <b>241</b>
Execution Plan Generation . . . . .	241
Parser . . . . .	243
Algebrizer . . . . .	243
Optimization . . . . .	244
Execution Plan Caching . . . . .	251
Components of the Execution Plan . . . . .	251
Query Plan . . . . .	251
Execution Context . . . . .	251
Aging of the Execution Plan . . . . .	252
Analyzing the Execution Plan Cache . . . . .	252
Execution Plan Reuse . . . . .	253
Ad Hoc Workload . . . . .	254
Prepared Workload . . . . .	255
Plan Reusability of an Ad Hoc Workload . . . . .	256
Plan Reusability of a Prepared Workload . . . . .	264



Query Plan Hash and Query Hash . . . . .	274
Execution Plan Cache Recommendations . . . . .	278
Explicitly Parameterize Variable Parts of a Query . . . . .	278
Create Stored Procedures to Implement	
Business Functionality . . . . .	278
Code with sp_executesql to Avoid Stored	
Procedure Maintenance . . . . .	278
Implement the Prepare/Execute Model to	
Avoid Resending a Query String . . . . .	279
Avoid Ad Hoc Queries . . . . .	279
Prefer sp_executesql over EXECUTE for Dynamic Queries . . . . .	279
Parameterize Variable Parts of Queries with Care . . . . .	280
Do Not Allow Implicit Resolution of Objects in Queries . . . . .	280
Summary . . . . .	281

## ■ CHAPTER 10    **Stored Procedure Recompilation** . . . . . 283

Benefits and Drawbacks of Recompilation . . . . .	283
Identifying the Statement Causing Recompilation . . . . .	286
Analyzing Causes of Recompilation . . . . .	288
Schema or Bindings Changes . . . . .	289
Statistics Changes . . . . .	289
Deferred Object Resolution . . . . .	292
SET Options Changes . . . . .	294
Execution Plan Aging . . . . .	295
Explicit Call to sp_recompile . . . . .	295
Explicit Use of the RECOMPILE Clause . . . . .	296
Avoiding Recompilations . . . . .	298
Do Not Interleave DDL and DML Statements . . . . .	298
Avoiding Recompilations Caused by Statistics Change . . . . .	300
Using Table Variables . . . . .	302
Avoiding Changing SET Options Within a Stored Procedure . . . . .	304
Using OPTIMIZE FOR Query Hint . . . . .	305
Using Plan Guides . . . . .	307
Summary . . . . .	311

<b>CHAPTER 11</b>	<b>Query Design Analysis</b>	313
	Query Design Recommendations	313
	Operating on Small Result Sets	314
	Limit the Number of Columns in select_list	314
	Use Highly Selective WHERE Clauses	315
	Using Indexes Effectively	316
	Avoid Nonsargable Search Conditions	316
	Avoid Arithmetic Operators on the WHERE Clause Column	320
	Avoid Functions on the WHERE Clause Column	322
	Avoiding Optimizer Hints	324
	JOIN Hint	325
	INDEX Hints	327
	Using Domain and Referential Integrity	329
	NOT NULL Constraint	329
	Declarative Referential Integrity	332
	Avoiding Resource-Intensive Queries	334
	Avoid Data Type Conversion	335
	Use EXISTS over COUNT(*) to Verify Data Existence	337
	Use UNION ALL Instead of UNION	338
	Use Indexes for Aggregate and Sort Conditions	339
	Avoid Local Variables in a Batch Query	339
	Be Careful Naming Stored Procedures	344
	Reducing the Number of Network Round-Trips	345
	Execute Multiple Queries Together	346
	Use SET NOCOUNT	346
	Reducing the Transaction Cost	346
	Reduce Logging Overhead	347
	Reduce Lock Overhead	348
	Summary	350
<b>CHAPTER 12</b>	<b>Blocking Analysis</b>	351
	Blocking Fundamentals	351
	Understanding Blocking	352
	Atomicity	353
	Consistency	355
	Isolation	356
	Durability	356

Database Locks .....	357
Lock Granularity .....	357
Lock Escalation .....	362
Lock Modes .....	362
Lock Compatibility .....	369
Isolation Levels .....	369
Read Uncommitted .....	370
Read Committed .....	371
Repeatable Read .....	372
Serializable .....	375
Snapshot .....	381
Effect of Indexes on Locking .....	381
Effect of a Nonclustered Index .....	382
Effect of a Clustered Index .....	384
Effect of Indexes on the Serializable Isolation Level .....	385
Capturing Blocking Information .....	385
Capturing Blocking Information with SQL .....	386
Profiler Trace and the Blocked Process Report Event .....	388
Blocking Resolutions .....	390
Optimize the Queries .....	390
Decrease the Isolation Level .....	391
Partition the Contended Data .....	392
Covering Index on Contended Data .....	392
Recommendations to Reduce Blocking .....	393
Automation to Detect and Collect Blocking Information .....	394
Summary .....	399

<b>CHAPTER 13</b>	<b>Deadlock Analysis .....</b>	<b>401</b>
	Deadlock Fundamentals .....	401
	Choosing the Deadlock Victim .....	402
	Using Error Handling to Catch a Deadlock .....	403
	Deadlock Analysis .....	403
	Collecting Deadlock Information .....	404
	Analyzing the Deadlock .....	405
	Avoiding Deadlocks .....	410
	Accessing Resources in the Same Chronological Order .....	410
	Decreasing the Number of Resources Accessed .....	411
	Minimizing Lock Contention .....	411
	Summary .....	413

<b>CHAPTER 14</b>	<b>Cursor Cost Analysis</b>	415
	Cursor Fundamentals	415
	Cursor Location	417
	Cursor Concurrency	418
	Cursor Types	419
	Cursor Cost Comparison	422
	Cost Comparison on Cursor Location	422
	Cost Comparison on Cursor Concurrency	424
	Cost Comparison on Cursor Type	426
	Default Result Set	428
	Benefits	429
	Drawbacks	429
	Analyzing SQL Server Overhead with Cursors	432
	Analyzing SQL Server Overhead with T-SQL Cursors	432
	Cursor Recommendations	436
	Summary	437
 <b>CHAPTER 15</b>	 <b>Database Workload Optimization</b>	 439
	Workload Optimization Fundamentals	439
	Workload Optimization Steps	440
	Sample Workload	441
	Capturing the Workload	444
	Analyzing the Workload	446
	Identifying the Costliest Query	447
	Determining the Baseline Resource Use of the Costliest Query	449
	Overall Resource Use	449
	Detailed Resource Use	449
	Analyzing and Optimizing External Factors	452
	Analyzing the Batch-Level Options Used by the Application	452
	Analyzing the Effectiveness of Statistics	453
	Analyzing the Need for Defragmentation	454
	Analyzing the Internal Behavior of the Costliest Query	458
	Analyzing the Query Execution Plan	458
	Identifying the Costly Steps in the Execution Plan	459
	Analyzing the Effectiveness of the Processing Strategy	460
	Optimizing the Costliest Query	460
	Modifying an Existing Index	461
	Analyzing the Application of a Join Hint	463
	Avoiding the Clustered Index Scan Operation	466
	Modifying the Procedure	466

Analyzing the Effect on Database Workload .....	469
Iterating Through Optimization Phases .....	471
Summary .....	473
 <b>CHAPTER 16 SQL Server Optimization Checklist</b> .....	 475
Database Design .....	475
Balancing Under- and Overnormalization .....	476
Benefiting from Entity-Integrity Constraints .....	477
Benefiting from Domain and Referential Integrity Constraints .....	479
Adopting Index-Design Best Practices .....	480
Avoiding the Use of the sp_ Prefix for Stored Procedure Names .....	482
Minimizing the Use of Triggers .....	482
Query Design .....	482
Use the Command SET NOCOUNT ON .....	483
Explicitly Define the Owner of an Object .....	483
Avoid Nonsargable Search Conditions .....	483
Avoid Arithmetic Operators on the WHERE Clause Column .....	484
Avoid Optimizer Hints .....	485
Stay Away from Nesting Views .....	485
Ensure No Implicit Data Type Conversions .....	485
Minimize Logging Overhead .....	485
Adopt Best Practices for Reusing Execution Plans .....	486
Adopt Best Practices for Database Transactions .....	487
Eliminate or Reduce the Overhead of Database Cursors .....	488
Configuration Settings .....	488
Affinity Mask .....	488
Memory Configuration Options .....	488
Cost Threshold for Parallelism .....	489
Max Degree of Parallelism .....	489
Optimize for Ad Hoc Workloads .....	489
Query Governor Cost Limit .....	490
Fill Factor (%) .....	490
Blocked Process Threshold .....	490
Database File Layout .....	490
Database Compression .....	491

Database Administration .....	491
Keep the Statistics Up-to-Date .....	491
Maintain a Minimum Amount of Index Defragmentation.....	492
Cycle the SQL Error Log File.....	492
Avoid Automatic Database Functions Such As	
AUTO_CLOSE or AUTO_SHRINK.....	492
Minimize the Overhead of SQL Tracing .....	493
Database Backup.....	493
Incremental and Transaction Log Backup Frequency .....	493
Backup Distribution .....	494
Backup Compression.....	495
Summary.....	495
 ■ INDEX .....	 497